

# bdf

Brighton Data Forum

connect    grow    learn

# from zero to query

a sql primer

oskar 2023-03-06

# sql - a fundamental tool for the data professional

- database management
- data pipeline engineering
- data modeling
- data designing
- **big** data (parallel, distributed)
- data querying
- data analytics

# sql - a fundamental tool for the data professional

- database management
- data pipeline engineering
- data modeling
- data designing
- big data (parallel, distributed)
- data querying
- data analytics

**transactional database**

**analytical database**

**schema**

**table**


**row**

**column**

**column**

**row**

**value**



data definition	data management	data querying	data control	transaction control
to operate on entire tables	to operate on table values, rows, columns	to fetch data from tables	to control access to schemas + tables	for transactional atomicity, dev
CREATE	INSERT	SELECT	GRANT	COMMIT
DROP	UPDATE		REVOKE	ROLLBACK
ALTER	DELETE			SAVE POINT
TRUNCATE				

data definition	data management	data querying	data control	transaction control
to operate on entire tables	to operate on table cells, rows, columns	to fetch data from tables	to control access to schemas + tables	for transactional atomicity, dev
CREATE	INSERT	SELECT	GRANT	COMMIT
DROP	UPDATE		REVOKE	ROLLBACK
ALTER	DELETE			SAVE POINT
TRUNCATE				

# a note on `sqlite`

- small (<2mb)
- open source
- serverless
- self-contained
- fast
- complete
- in-memory
- cross-platform
- ubiquitous





# sqlite commands



- these are not sql commands!
- they start with a '.'
- they operate on the environment, not the data
- examples:
  - .quit
  - .open <path-to-database>
  - .show
  - .help
  - .cd <directory>
  - .shell CMD ARGS...

```
.open data/sqlite-sakila.db
```

```
.header ON
```

```
.mode qbox
```

```
.show
```

```
.tables
```

# .tables

```
sqlite> .tables
```

actor

address

category

city

country

customer

customer\_list

```
sqlite> █
```

film

film\_actor

film\_category

film\_list

film\_text

inventory

language

payment

rental

sales\_by\_film\_category

sales\_by\_store

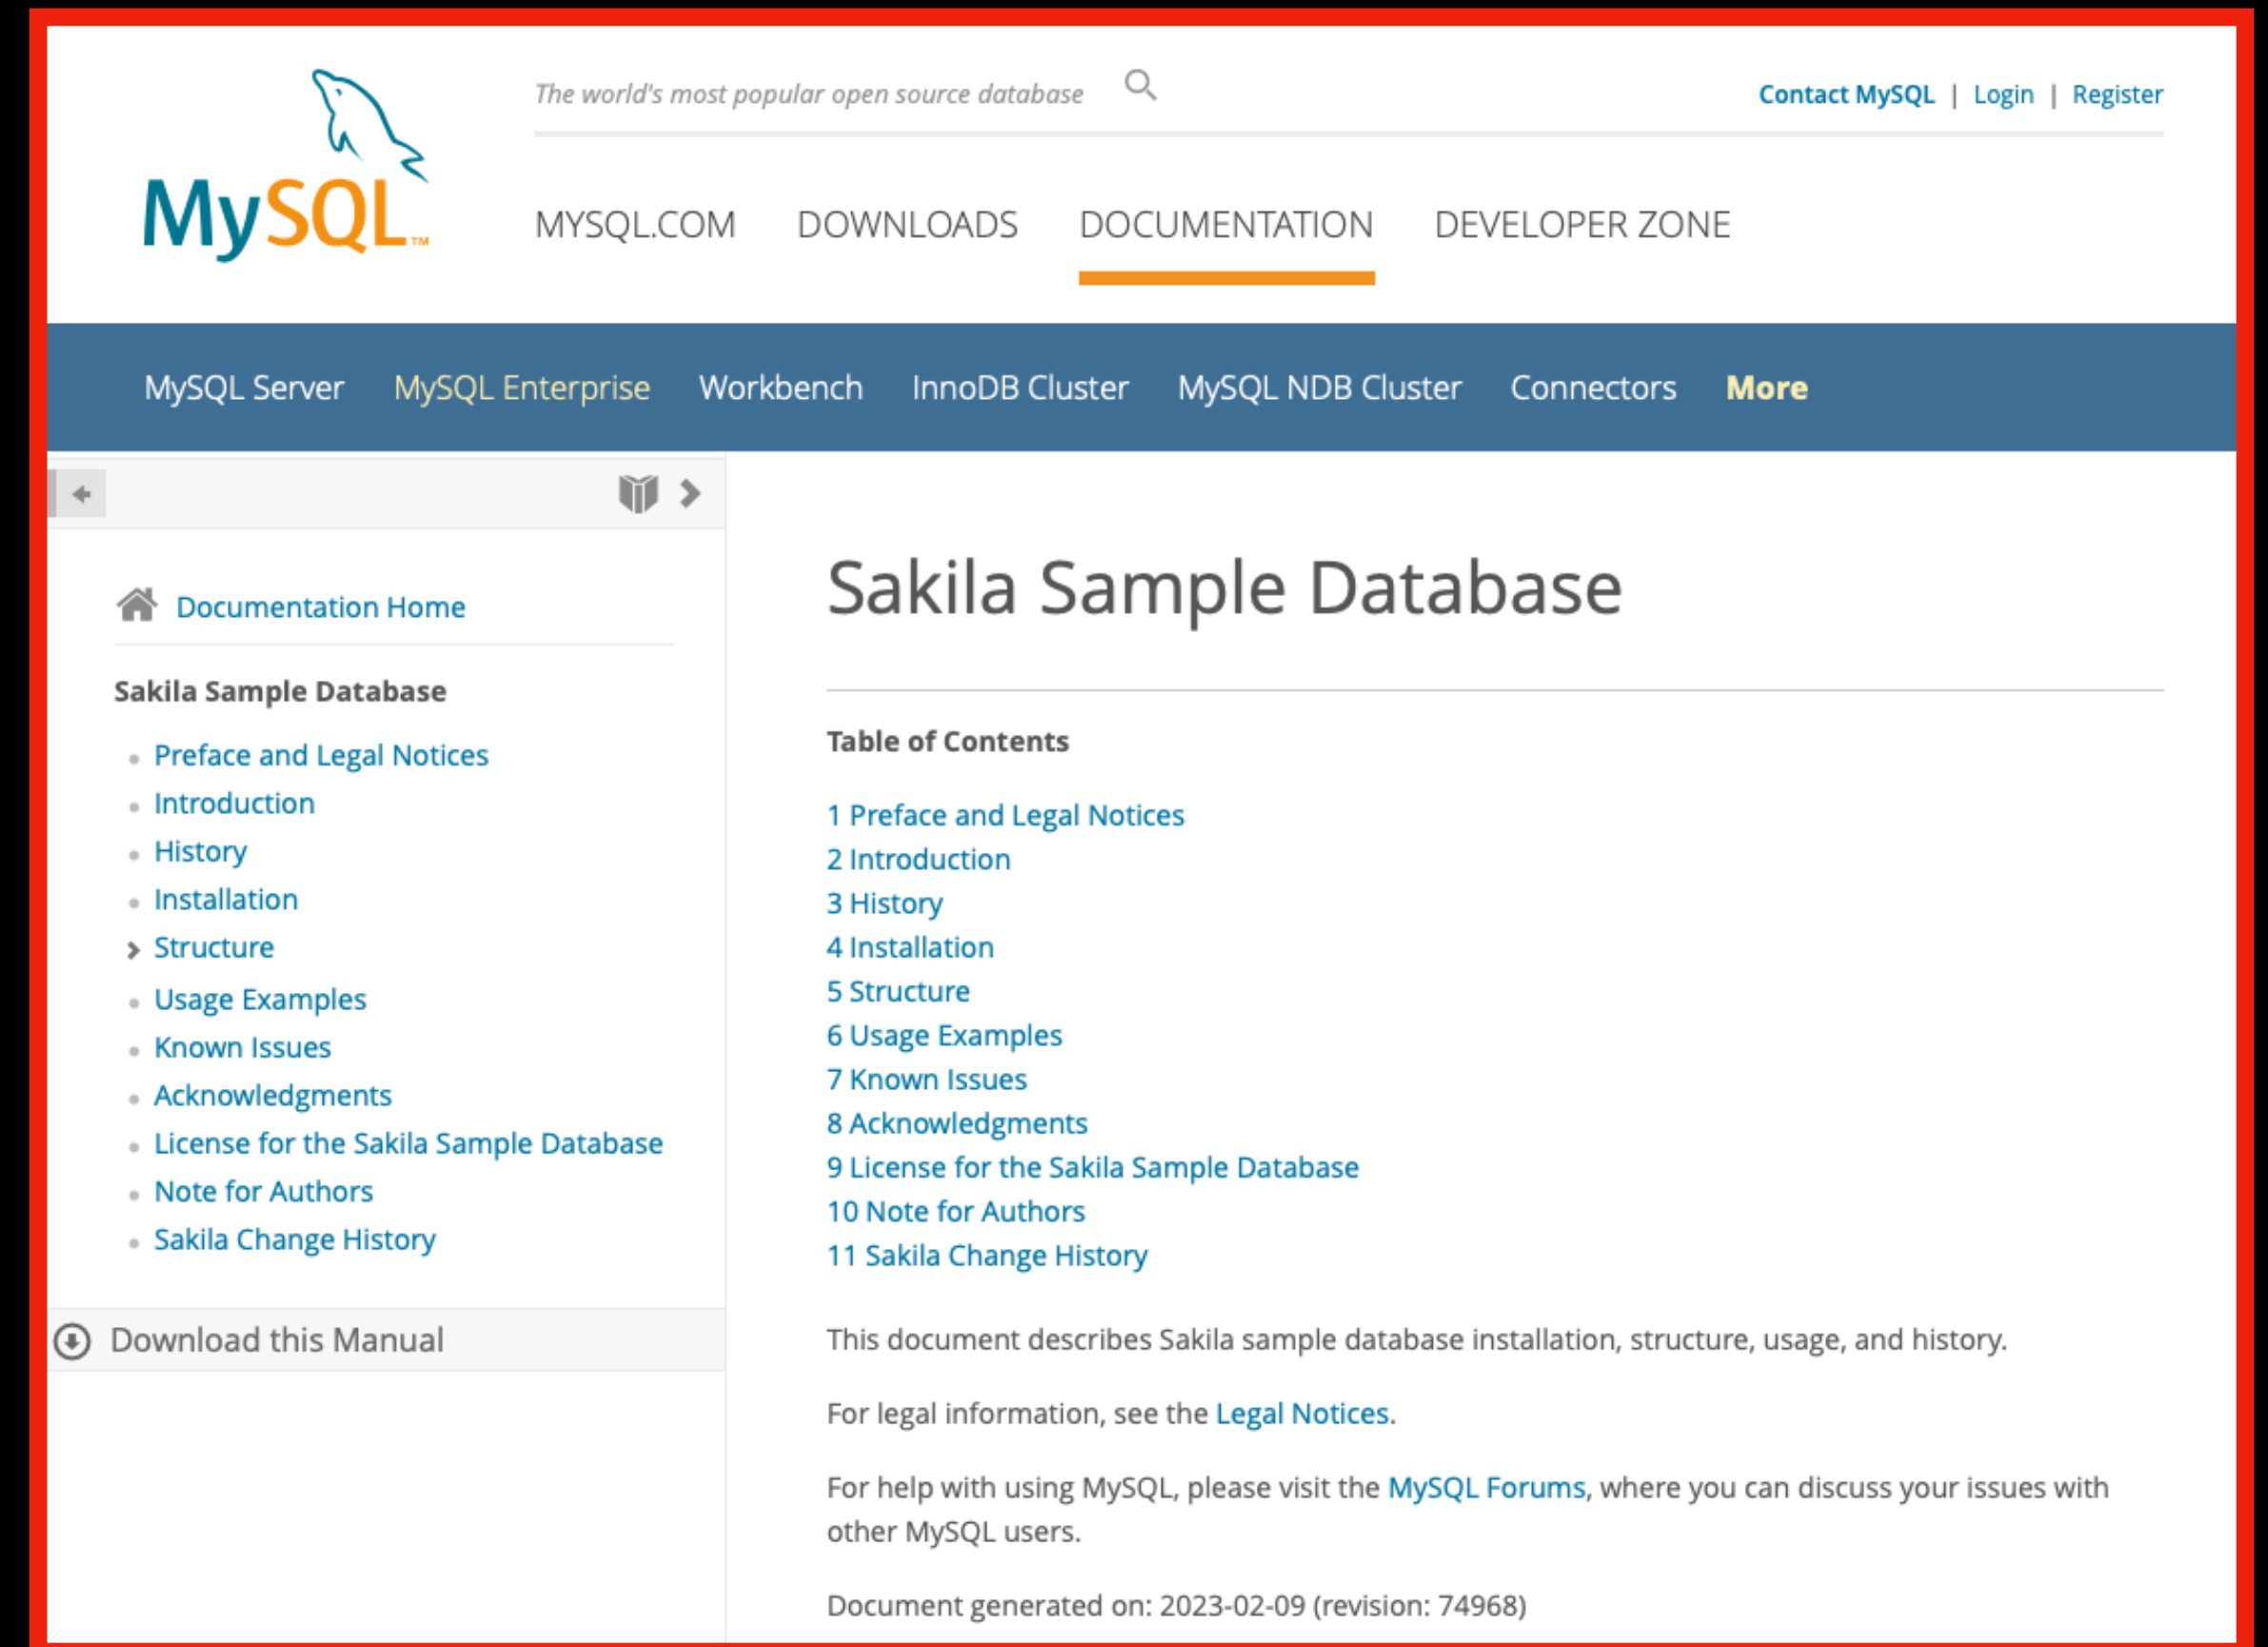
staff

staff\_list

store

# the sakila training data

- classic, fictional data
- dvd rental company
- 20 relational tables:
  - normalised: no repetition
  - stores
  - inventory
  - films
  - film casting
  - actors
  - film ratings



The screenshot shows the MySQL website's documentation page for the Sakila Sample Database. The page is titled "Sakila Sample Database" and features a "Table of Contents" with 11 items: 1 Preface and Legal Notices, 2 Introduction, 3 History, 4 Installation, 5 Structure, 6 Usage Examples, 7 Known Issues, 8 Acknowledgments, 9 License for the Sakila Sample Database, 10 Note for Authors, and 11 Sakila Change History. The page also includes a "Download this Manual" button and a footer with the date "2023-02-09 (revision: 74968)".

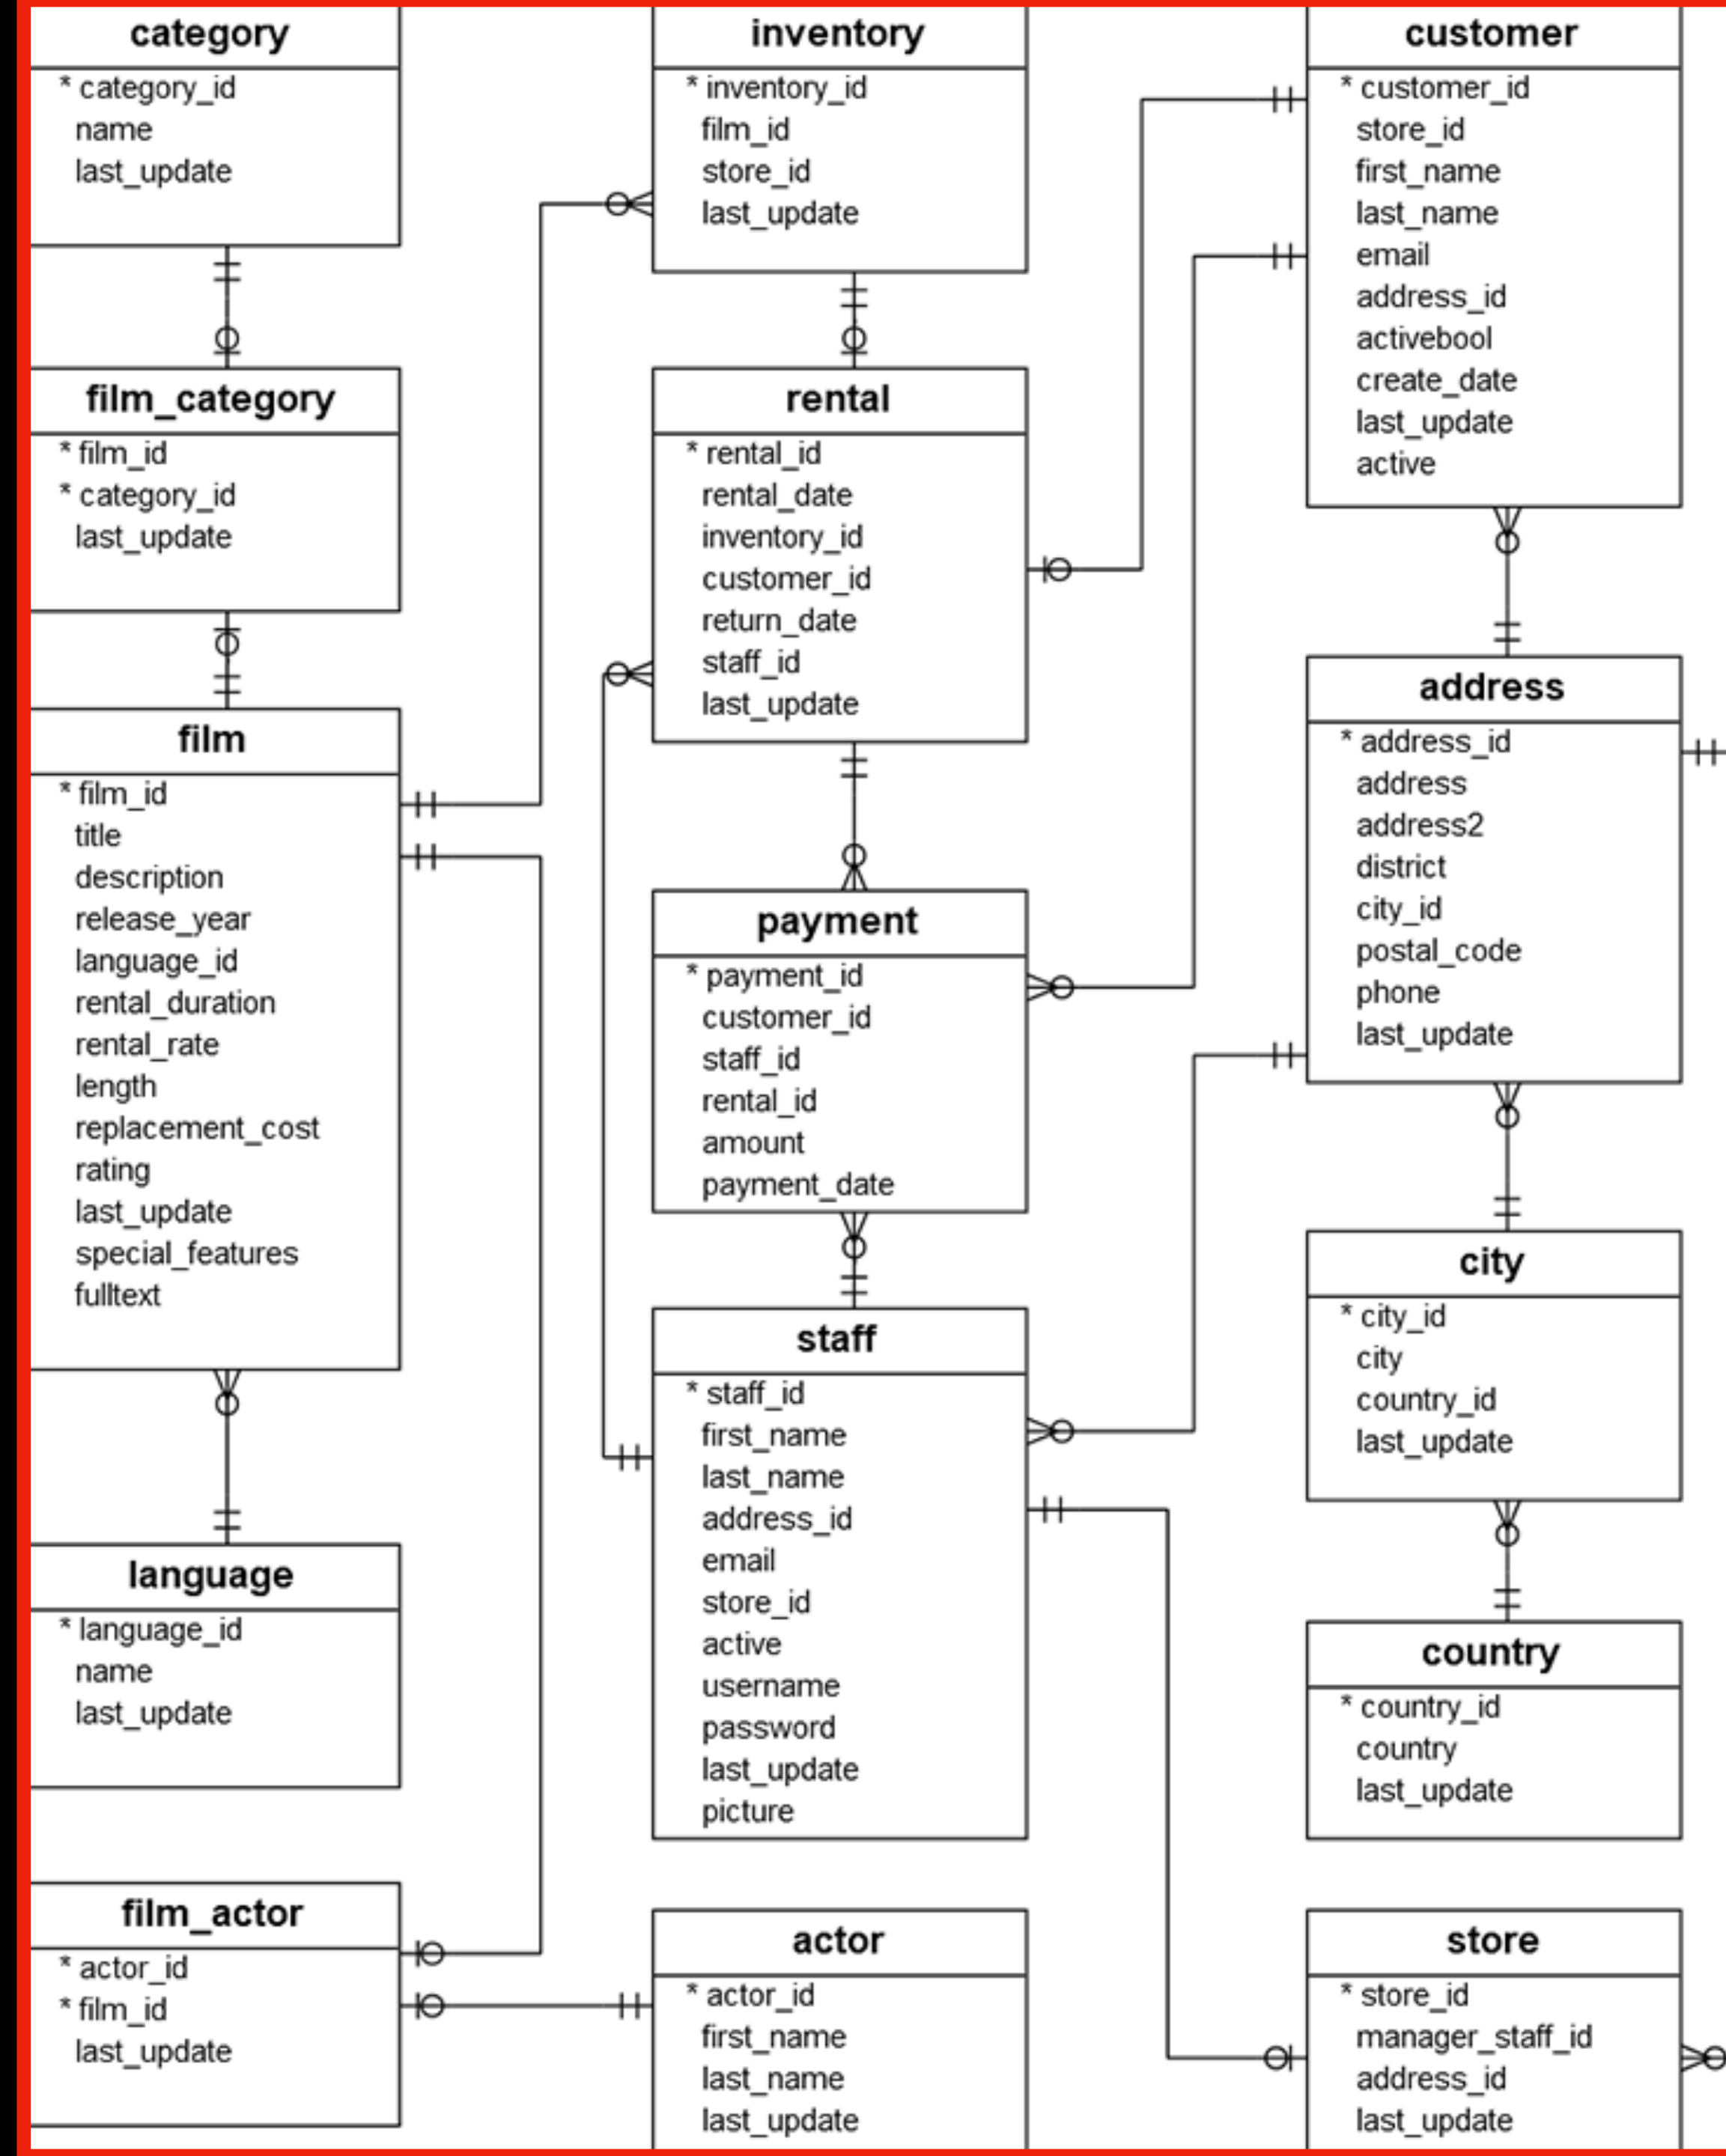
The MySQL logo is visible in the top left corner, and the tagline "The world's most popular open source database" is in the top right. The navigation bar includes links to MySQL.COM, DOWNLOADS, DOCUMENTATION (highlighted), and DEVELOPER ZONE. Below the navigation bar, there are links to MySQL Server, MySQL Enterprise, Workbench, InnoDB Cluster, MySQL NDB Cluster, Connectors, and More.

The left sidebar contains a "Documentation Home" link and a "Sakila Sample Database" section with a list of links: Preface and Legal Notices, Introduction, History, Installation, Structure, Usage Examples, Known Issues, Acknowledgments, License for the Sakila Sample Database, Note for Authors, and Sakila Change History.

The main content area is titled "Sakila Sample Database" and contains a "Table of Contents" section with a list of 11 items. Below the table of contents, there is a paragraph stating: "This document describes Sakila sample database installation, structure, usage, and history. For legal information, see the [Legal Notices](#). For help with using MySQL, please visit the [MySQL Forums](#), where you can discuss your issues with other MySQL users." The footer of the page states: "Document generated on: 2023-02-09 (revision: 74968)".

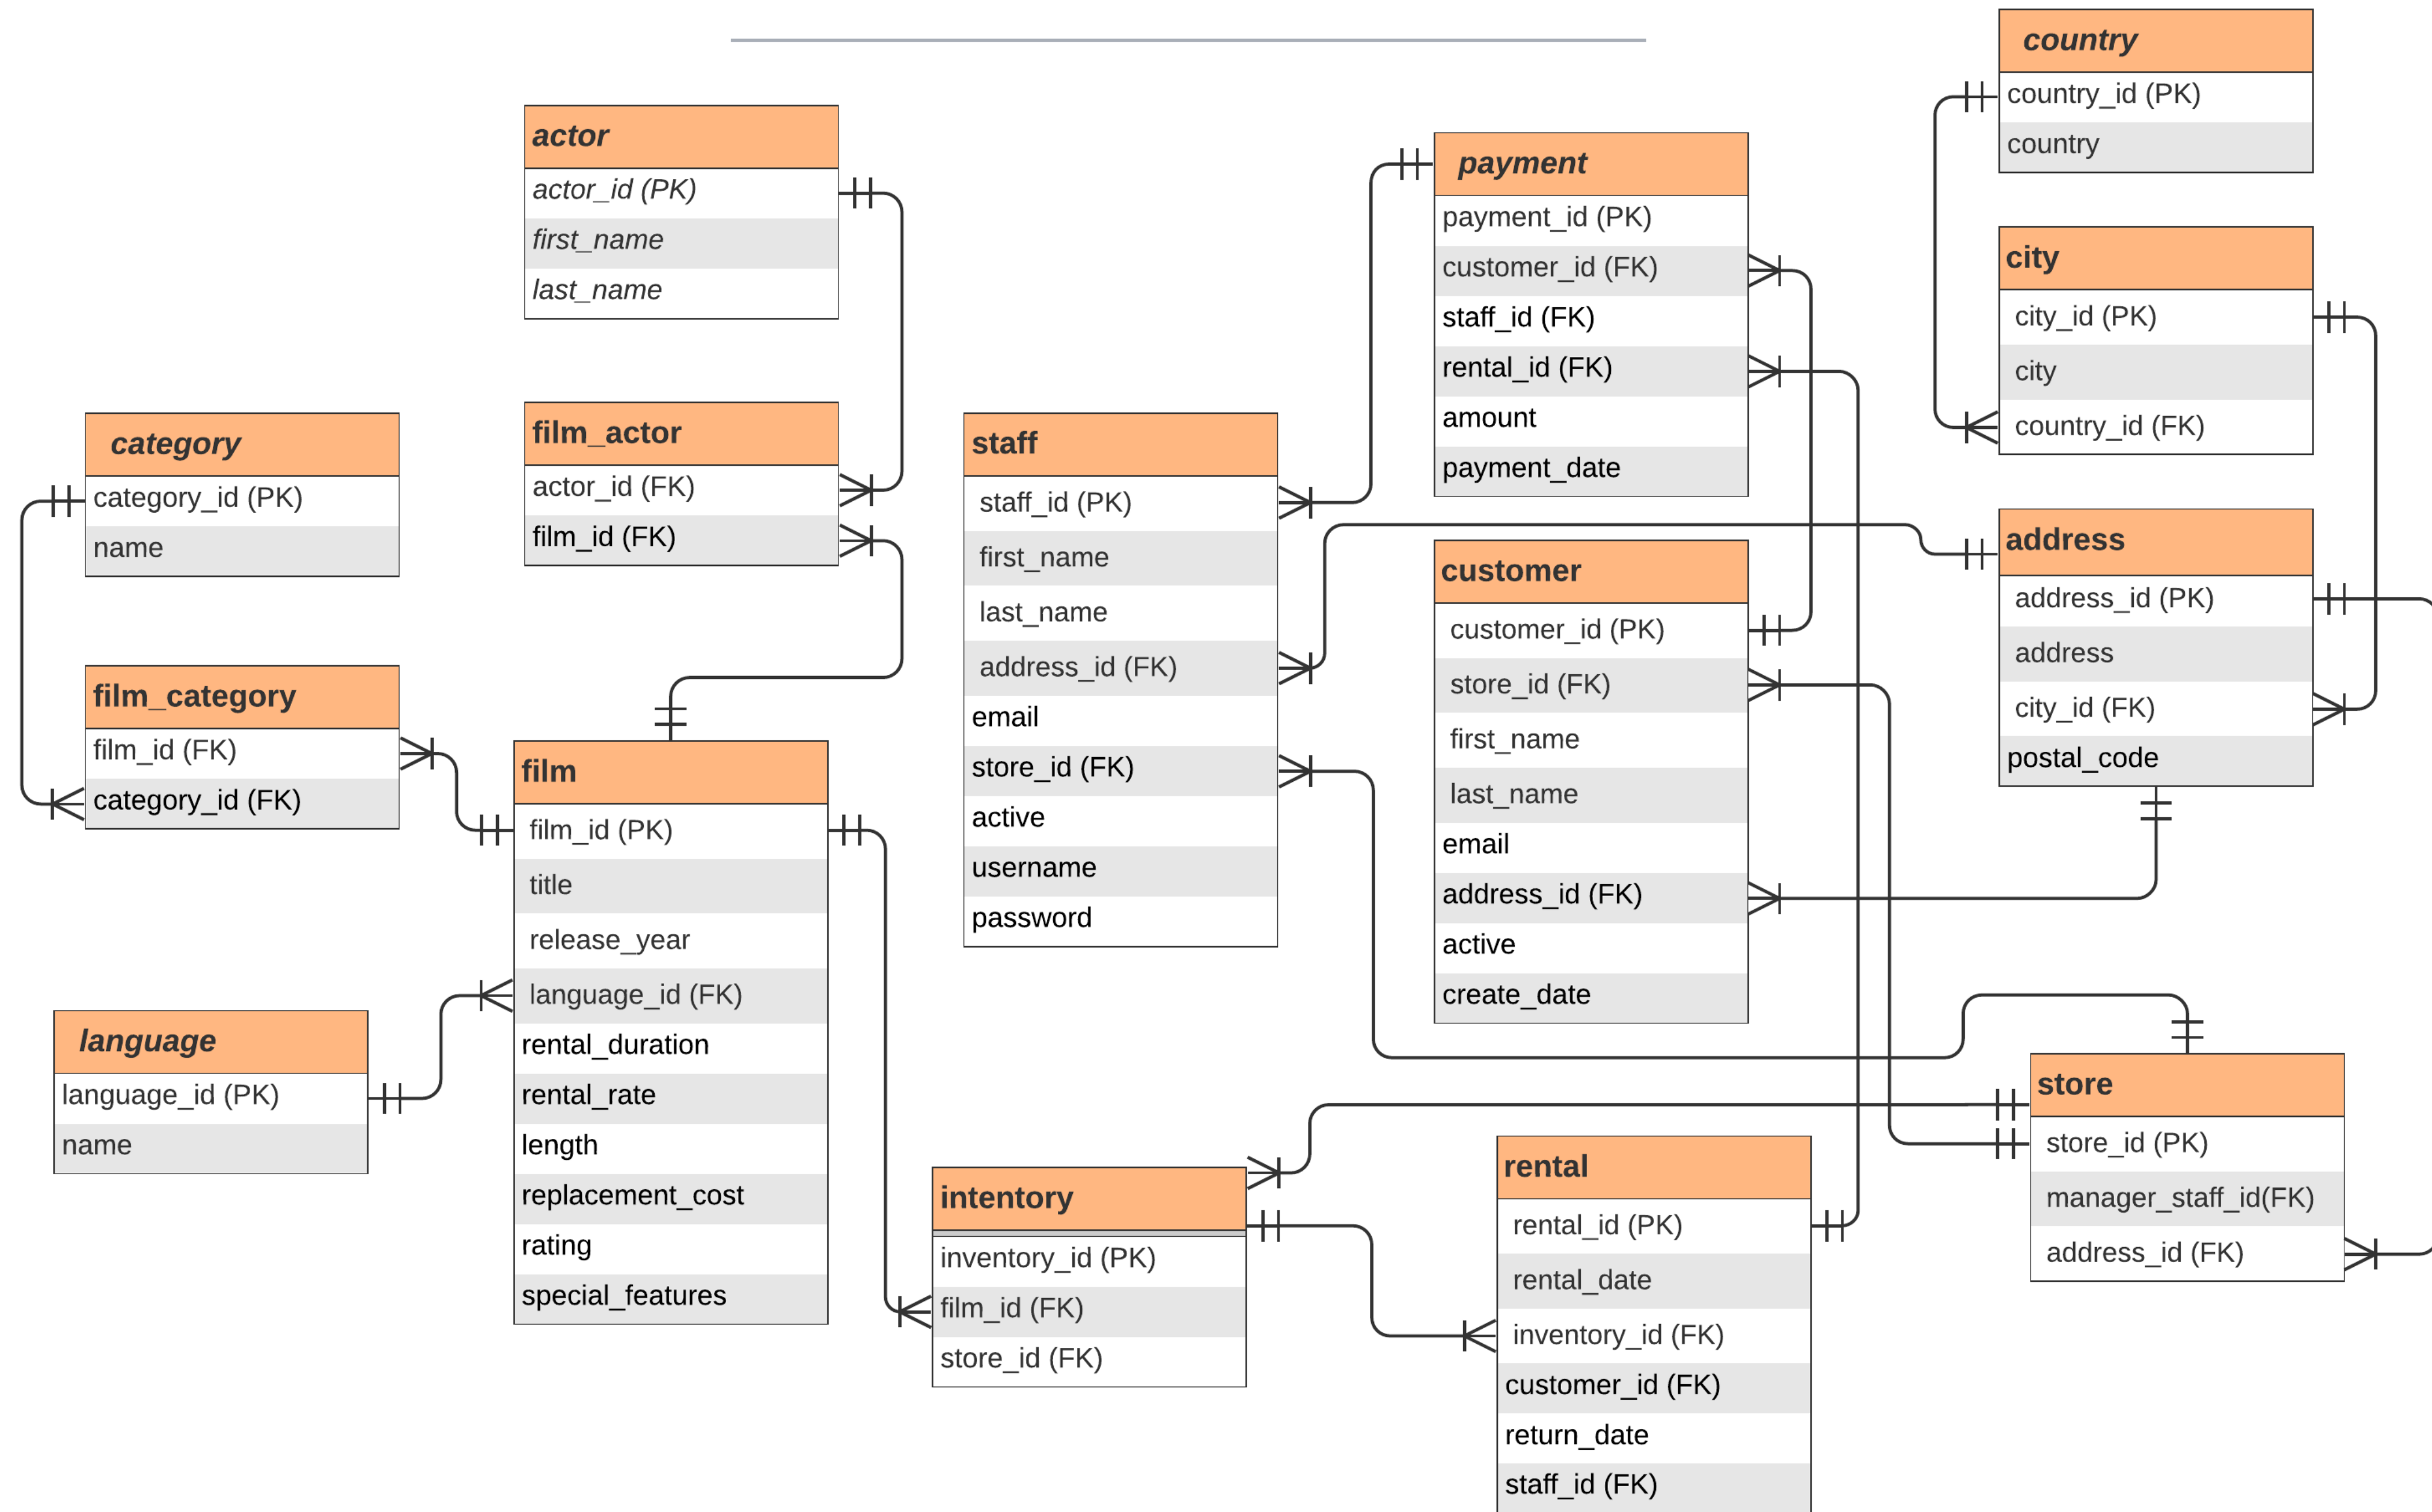
# the sakila training data

- classic, fictional data
- dvd rental company
- 20 relational tables:
  - normalised: no repetition
  - stores
  - inventory
  - films
  - film casting
  - actors
  - film ratings






SQLite3 Sakila Sample Database ERD



today's objective:

“which **top 10** actors were rented out the greatest number of times, counting only ‘**R**’ rated films made in **2006**?”



# sql commands

- these run on the database
- they end with a ';' ;
- you can add comments with '-- a comment'
- they operate on the data tables
- example:
  - SELECT {columns} FROM table; -- a&b



# today's plan:

“which top 10 actors were rented out the greatest number of times, counting only ‘R’ rated films made in 2006?”

- SELECT {columns} FROM {table};
- + LIMIT num
- + WHERE {a\_condition}
- + ORDER BY {columns}
- + GROUP BY {columns}
- + HAVING {a\_condition}
- + INNER JOIN {table\_2} ON {col1}={col2}

# what do the tables contain?

“which top 10 actors were rented out the greatest number of times, counting only ‘R’ rated films made in 2006?”

- SELECT {columns} FROM {table};
- + LIMIT num
- + WHERE {a\_condition}
- + ORDER BY {columns}
- + GROUP BY {columns}
- + HAVING {a\_condition}
- + [INNER] JOIN {table\_2} ON {col1}={col2}

# SELECT ... FROM ...;

- `SELECT * FROM {tablename};`  
-- returns all columns and all rows from a table
- `SELECT name, category_id FROM category;`  
-- returns columns name and id (in that order) of each category
- `SELECT a.first_name AS name FROM actor a;`  
-- creates an alias for table actor, fetches column first\_name as name
- `SELECT rental_rate + replacement_cost AS total_cost FROM film;`  
-- returns the 'total\_cost' of renting, then replacing a film
- `SELECT DISTINCT first_name FROM actor;`  
-- returns all the first names in the actor table, no duplicates

# SELECT (aggregate function) FROM ... ;

- SELECT COUNT(\*) AS num\_records FROM actor;  
-- returns the number of rows in table actor, names the output 'num\_records'
- SELECT COUNT(DISTINCT rating) FROM film;  
-- returns a count of distinct values in the rating column
- SELECT AVG(replacement\_cost) AS avg\_cost FROM film;  
-- returns the average replacement cost of a film
- SELECT AVG(rental\_rate) AS average\_rental\_rate FROM film;  
-- returns the average rate of rental from film table
- SELECT MAX(s.sale\_cost) AS highest\_value\_sale FROM sales s;  
-- returns the highest value sale from sales
- SELECT MIN(length) AS shortest\_length FROM film;  
-- returns the length of the shortest film

# that's too many rows!

“which top 10 actors were rented out the greatest number of times, counting only ‘R’ rated films made in 2006?”

- SELECT {columns} FROM {table}
- + LIMIT num;
- + WHERE {a\_condition}
- + ORDER BY {columns}
- + GROUP BY {columns}
- + HAVING {a\_condition}
- + [INNER] JOIN {table\_2} ON {col1}={col2}

# SELECT ... FROM ... LIMIT ...;

- `SELECT * FROM {table} LIMIT {n};`  
-- returns {n} unspecified rows of all columns from {table}
- `SELECT * FROM sales LIMIT 5;`  
-- returns 5 unspecified rows of all columns from sales
- `SELECT title, release_year FROM film LIMIT 15;`  
-- returns 15 unspecified rows of two columns from sales table
- `SELECT rental_id, rental_date FROM rental LIMIT 10;`  
-- returns region id and region name for 10 unspecified rows

but i only want specific rows!

“which top 10 actors were rented out the greatest number of times, counting only ‘R’ rated films made in 2006?”

- SELECT {columns} FROM {table}
- + LIMIT num
- + WHERE {a\_condition};
- + ORDER BY {columns}
- + GROUP BY {columns}
- + HAVING {a\_condition}
- + [INNER] JOIN {table\_2} ON {col1}={col2}

# SELECT ... FROM ... WHERE ... [LIMIT n];

- **SELECT \* FROM {table} WHERE {column}={expression};**  
-- returns only rows where the value in {column} equals {expression}
- **SELECT \* FROM table\_name WHERE column1<>{expression};**  
-- returns only rows where the value in column1 is not {expression}
- **SELECT title AS film\_name, rental\_rate FROM film WHERE rental\_rate<=1.0;**  
-- returns titles of films whose rental price is at most £1
- **SELECT name FROM items WHERE item\_price>=10 LIMIT 8;**  
-- returns 8 of the items whose price is greater or equal to £10



# comparison operators

operator syntax	meaning
{column} = {expression}	column value is equal to expression value
{column} <> {expression}	column value is not equal to expression value
{column} != {expression}	column value is not equal to expression value
{column} < {expression}	column value is less than expression value
{column} <= {expression}	column value is less than or equal to expression value
{column} > {expression}	column value is greater than expression value
{column} >= {expression}	column value is greater than or equal to expression value
{column} IN ({exp1}, {exp2}, ...)	column value is one of 'exp1', 'exp2', ...
{column} LIKE '%expr%'	(string) column contains substring 'expr'
{column} BETWEEN {exp1} AND {exp2}	{exp1} <= column value <= {exp2}

# comparison operators

operator syntax	meaning
{column} = {expression}	column value is equal to expression value
{column} <> {expression}	column value is not equal to expression value
{column} != {expression}	column value is not equal to expression value
{column} < {expression}	column value is less than expression value
{column} <= {expression}	column value is less than or equal to expression value
{column} > {expression}	column value is greater than expression value
{column} >= {expression}	column value is greater than or equal to expression value
{column} IN ({exp1}, {exp2}, ...)	column value is one of 'exp1', 'exp2', ...
{column} LIKE '%expr%'	(string) column contains substring 'expr'
{column} BETWEEN {exp1} AND {exp2}	{exp1} <= column value <= {exp2}

# SELECT ... FROM ... WHERE ...;

- `SELECT * FROM rental WHERE rental_date BETWEEN '2005-11-01' AND '2005-01-01';`  
-- returns only rentals occurring in dece 2005
- `SELECT * FROM sales WHERE region_id IN (14,56,43) ;`  
-- returns only sales in regions with id 14, 56, or 43
- `SELECT * FROM region WHERE region_name LIKE '%new%';`  
-- returns only regions whose name contains 'new'
- `SELECT DISTINCT postal_code FROM address WHERE postal_code LIKE '97%';`  
-- show all the postal codes that start with '97'

but i only want the most extreme rows!

“which top 10 actors were rented out the greatest number of times, counting only ‘R’ rated films made in 2006?”

- SELECT {columns} FROM {table}
- + LIMIT num;
- + WHERE {a\_condition}
- + ORDER BY {columns}
- + GROUP BY {columns}
- + HAVING {a\_condition}
- + [INNER] JOIN {table\_2} ON {col1}={col2}

SELECT ... FROM ... ORDER BY ... LIMIT ...;

- SELECT \* FROM payment ORDER BY payment\_date LIMIT 7;  
-- return the earliest 7 payments in the payment table
- SELECT \* FROM payment ORDER BY payment\_date DESC LIMIT 7;  
-- return the latest 7 payments in the payment table
- SELECT \* FROM payment ORDER BY amount DESC LIMIT 5;;  
-- return only the top 5 highest payment amounts from the payment table

# how can i aggregate groups of rows into a single row?

“which top 10 actors were rented out the greatest number of times, counting only ‘R’ rated films made in 2006?”

- SELECT {columns} FROM {table}
- + LIMIT num
- + WHERE {a\_condition}
- + ORDER BY {columns}
- + GROUP BY {columns};
- + HAVING {a\_condition}
- + INNER JOIN {table\_2} ON {col1}={col2}

SELECT {col}, ... FROM ... GROUP BY {col};

- SELECT city\_id, COUNT(\*) AS num\_address FROM address GROUP BY city\_id;  
-- return number of addresses in each city in address table
- SELECT rating, AVG(length) AS avg\_len FROM film GROUP BY rating ORDER BY avg\_len;  
-- returns the average length of a movie in each rating category
- SELECT country\_id, COUNT(\*) AS num\_cities  
FROM city  
GROUP BY country\_id  
ORDER BY num\_cities DESC  
LIMIT 5;  
-- return top 5 country ids, by number of cities assigned to each

# how do i report only some aggregated groups?

“which top 10 actors were rented out the greatest number of times, counting only ‘R’ rated films made in 2006?”

- SELECT {columns} FROM {table};
- + LIMIT num
- + WHERE {a\_condition}
- + ORDER BY {columns}
- + GROUP BY {columns}
- + HAVING {a\_condition}
- + [INNER] JOIN {table\_2} ON {col1}={col2}



# SELECT ... FROM ... GROUP BY ... HAVING ...;

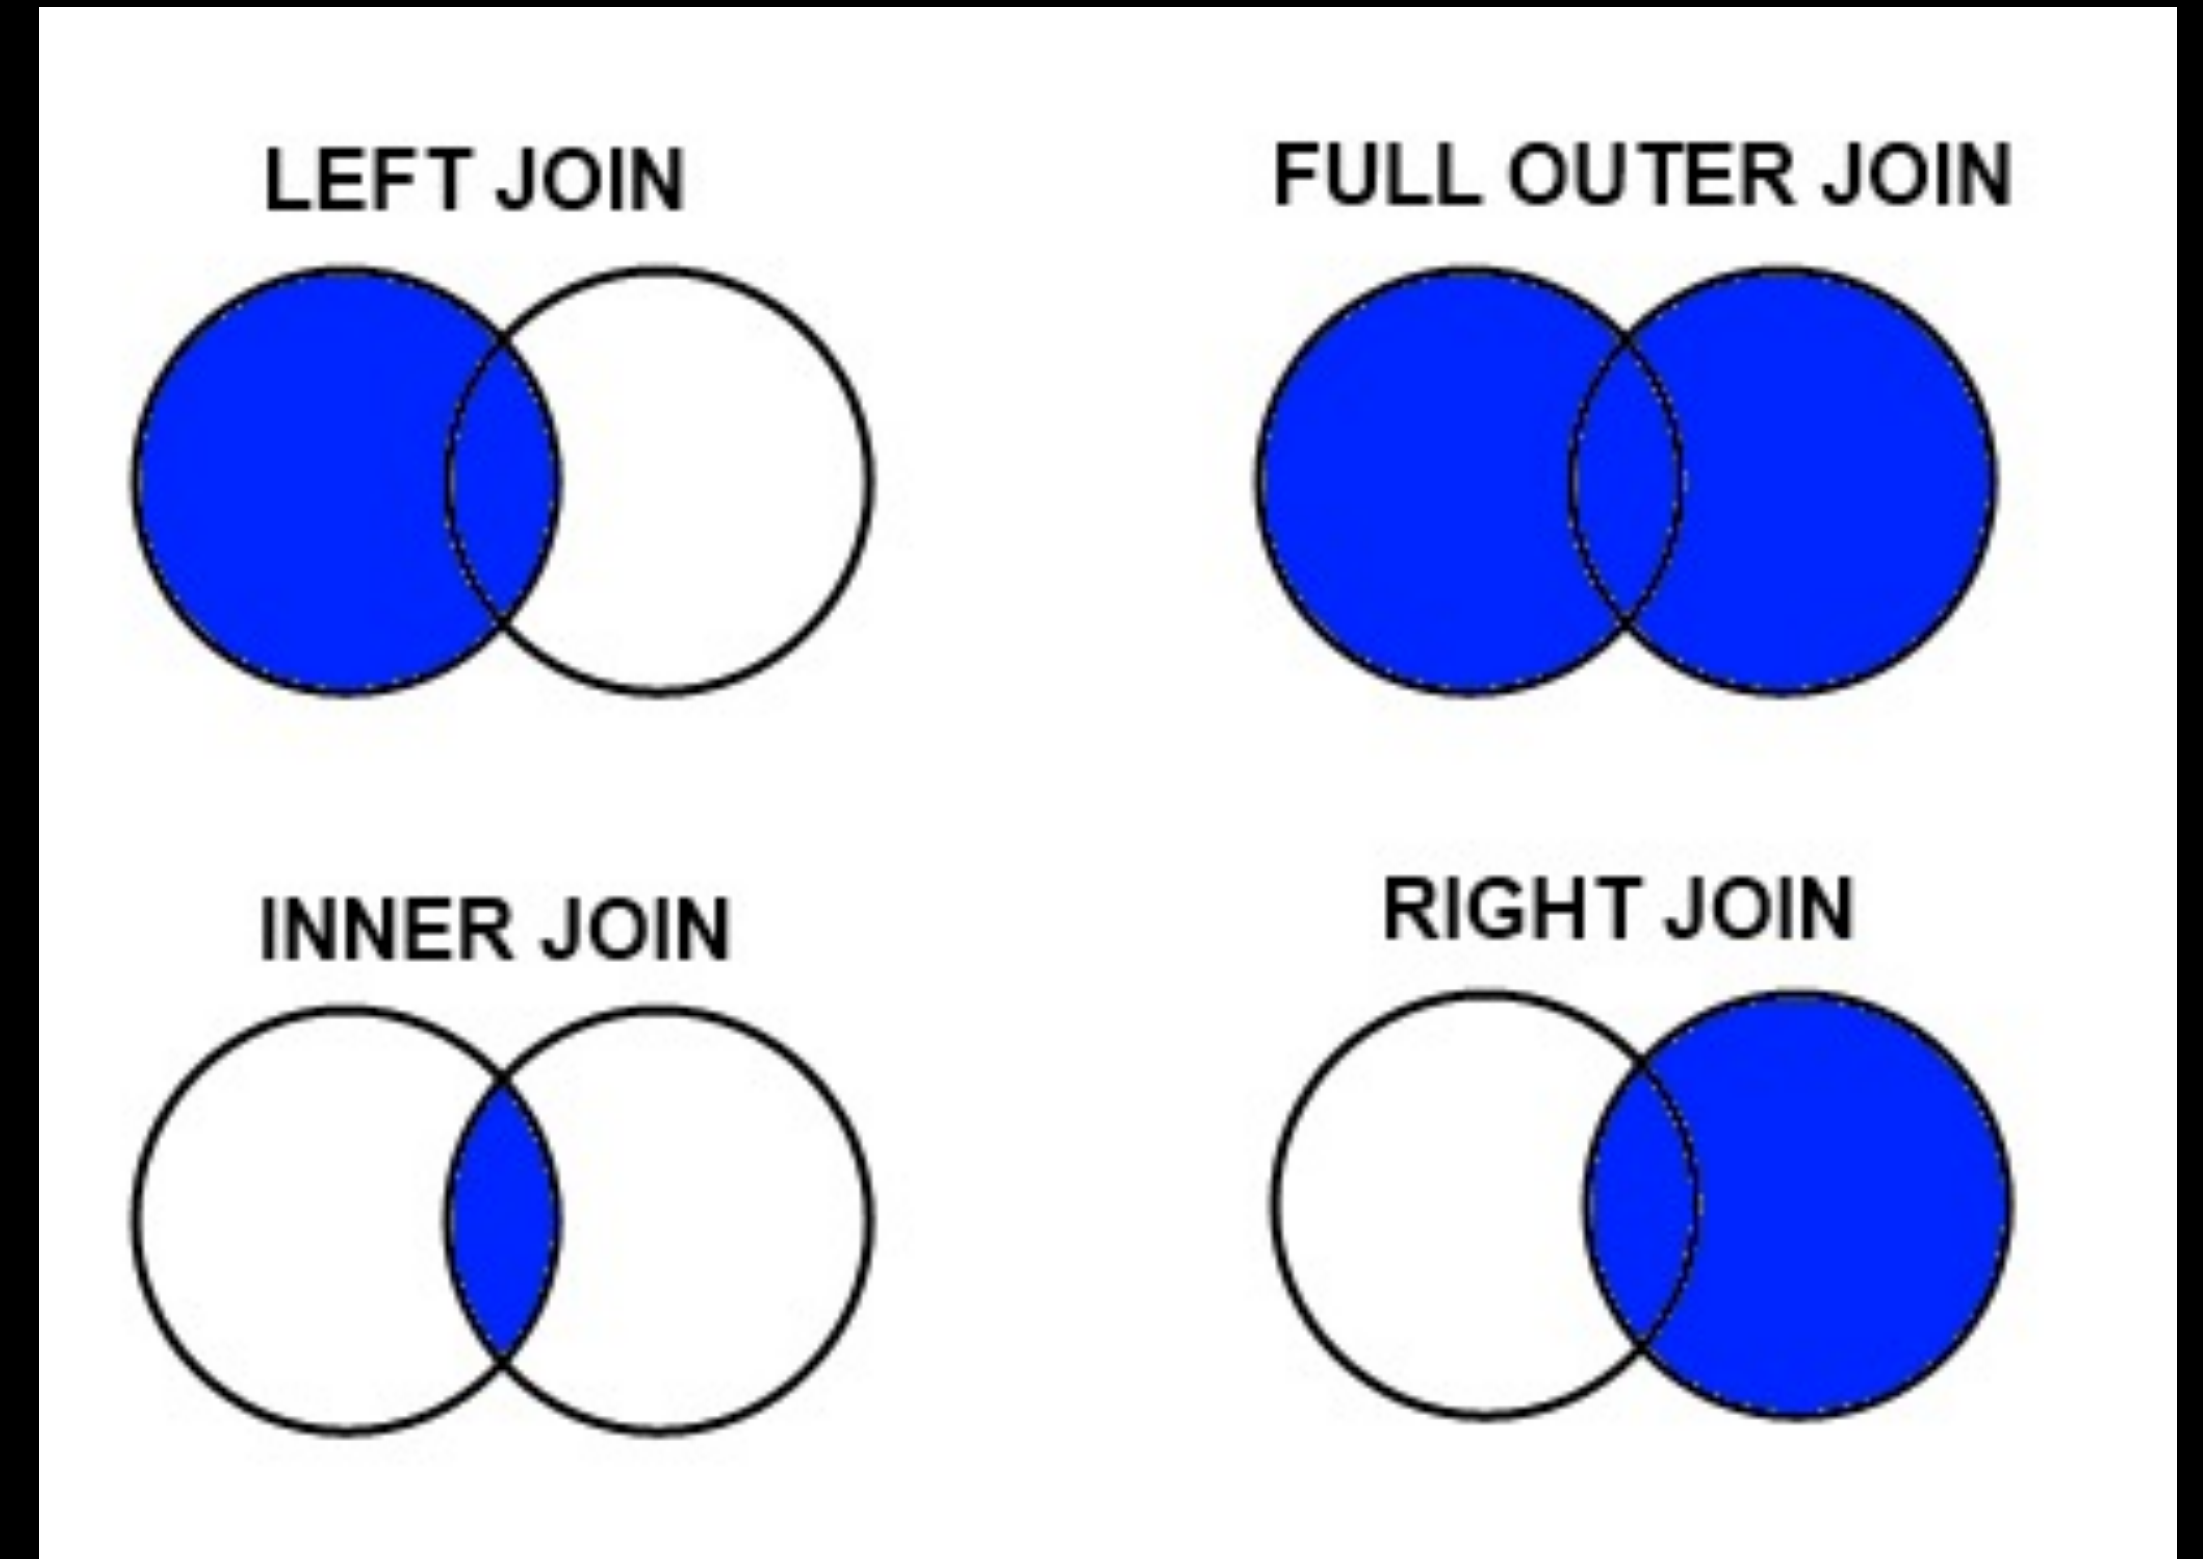
- `SELECT col1, COUNT(*) AS num FROM table GROUP BY col1 HAVING num>9;`  
-- count instances of each value of col1, but only output rows with count>9
- `SELECT rating, AVG(length) AS len FROM film GROUP BY rating HAVING len<115;`  
-- the film rating categories with average length of film under 115 minutes
- `SELECT actor_id, COUNT(*) AS n FROM film_actor GROUP BY actor_id HAVING n<15;`  
-- which actor ids have appeared in fewer than 15 films?

but my information is spread over two tables!

“which top 10 actors were rented out the greatest number of times, counting only ‘R’ rated films made in 2006?”

- SELECT {columns} FROM {table};
- + LIMIT num
- + WHERE {a\_condition}
- + ORDER BY {columns}
- + GROUP BY {columns}
- + HAVING {a\_condition}
- + [INNER] JOIN {table\_2} ON {col1}={col2}







# JOIN

city

city_id	city	country_id	last_update
1	A Corua (La Corua)	87	2021-03-06 15:51:49
2	Abha	82	2021-03-06 15:51:49
3	Abu Dhabi	101	2021-03-06 15:51:49
4	Acua	60	2021-03-06 15:51:49
5	Adana	97	2021-03-06 15:51:49
6	Addis Abeba	31	2021-03-06 15:51:49
7	Aden	107	2021-03-06 15:51:49
8	Adoni	44	2021-03-06 15:51:49

country

country_id	country	last_update
1	Afghanistan	2021-03-06 15:51:49
2	Algeria	2021-03-06 15:51:49
3	American Samoa	2021-03-06 15:51:49
4	Angola	2021-03-06 15:51:49
5	Anguilla	2021-03-06 15:51:49
6	Argentina	2021-03-06 15:51:49
7	Armenia	2021-03-06 15:51:49
8	Australia	2021-03-06 15:51:49
9	Austria	2021-03-06 15:51:49

# JOIN

city

city_id	city	country_id	last_update
1	A Corua (La Corua)	87	2021-03-06 15:51:49
2	Abha	82	2021-03-06 15:51:49
3	Abu Dhabi	101	2021-03-06 15:51:49
4	Acua	60	2021-03-06 15:51:49
5	Adana	97	2021-03-06 15:51:49
6	Addis Abeba	31	2021-03-06 15:51:49
7	Aden	107	2021-03-06 15:51:49
8	Adoni	44	2021-03-06 15:51:49

country

country_id	country	last_update
1	Afghanistan	2021-03-06 15:51:49
2	Algeria	2021-03-06 15:51:49
3	American Samoa	2021-03-06 15:51:49
4	Angola	2021-03-06 15:51:49
5	Anguilla	2021-03-06 15:51:49
6	Argentina	2021-03-06 15:51:49
7	Armenia	2021-03-06 15:51:49
8	Australia	2021-03-06 15:51:49
9	Austria	2021-03-06 15:51:49

# JOIN

city

city_id	city	country_id	last_update
1	A Corua (La Corua)	87	2021-03-06 15:51:49
2	Abha	82	2021-03-06 15:51:49
3	Abu Dhabi	101	2021-03-06 15:51:49
4	Acua	60	2021-03-06 15:51:49
5	Adana	97	2021-03-06 15:51:49
6	Addis Abeba	31	2021-03-06 15:51:49
7	Aden	107	2021-03-06 15:51:49
8	Adoni	44	2021-03-06 15:51:49

country

country_id	country	last_update
1	Afghanistan	2021-03-06 15:51:49
2	Algeria	2021-03-06 15:51:49
3	American Samoa	2021-03-06 15:51:49
4	Angola	2021-03-06 15:51:49
5	Anguilla	2021-03-06 15:51:49
6	Argentina	2021-03-06 15:51:49
7	Armenia	2021-03-06 15:51:49
8	Australia	2021-03-06 15:51:49
9	Austria	2021-03-06 15:51:49

# JOIN

city

city_id	city	country_id	last_update
1	A Corua (La Corua)	87	2021-03-06 15:51:49
2	Abha	82	2021-03-06 15:51:49
3	Abu Dhabi	101	2021-03-06 15:51:49
4	Acua	60	2021-03-06 15:51:49
5	Adana	97	2021-03-06 15:51:49
6	Addis Abeba	31	2021-03-06 15:51:49
7	Aden	107	2021-03-06 15:51:49
8	Adoni	44	2021-03-06 15:51:49

Turkey

country

country_id	country	last_update
1	Afghanistan	2021-03-06 15:51:49
2	Algeria	2021-03-06 15:51:49
3	American Samoa	2021-03-06 15:51:49
4	Angola	2021-03-06 15:51:49
5	Anguilla	2021-03-06 15:51:49
6	Argentina	2021-03-06 15:51:49
7	Armenia	2021-03-06 15:51:49
8	Australia	2021-03-06 15:51:49
9	Austria	2021-03-06 15:51:49

# JOIN

city

city_id	city	country_id	last_update
1	A Corua (La Corua)	87	2021-03-06 15:51:49
2	Abha	82	2021-03-06 15:51:49
3	Abu Dhabi	101	2021-03-06 15:51:49
4	Acua	60	2021-03-06 15:51:49
5	Adana	97	2021-03-06 15:51:49
6	Addis Abeba	31	2021-03-06 15:51:49
7	Aden	107	2021-03-06 15:51:49
8	Adoni	44	2021-03-06 15:51:49

Turkey

Yemen

country

country_id	country	last_update
1	Afghanistan	2021-03-06 15:51:49
2	Algeria	2021-03-06 15:51:49
3	American Samoa	2021-03-06 15:51:49
4	Angola	2021-03-06 15:51:49
5	Anguilla	2021-03-06 15:51:49
6	Argentina	2021-03-06 15:51:49
7	Armenia	2021-03-06 15:51:49
8	Australia	2021-03-06 15:51:49
9	Austria	2021-03-06 15:51:49



we want this

city-and-country

city_id	city	country
1	?	?
2	?	?
3	?	?
4	?	?
5	?	?
6	?	?
7	?	?
8	?	?

we want this


city-and-country

city_id	city	country
1	A Corua (La Corua)	Spain
2	Abha	Saudi Arabia
3	Abu Dhabi	United Arab Emirates
4	Acua	Mexico
5	Adana	Turkey
6	Addis Abeba	Ethiopia
7	Aden	Yemen
8	Adoni	India

so we add a JOIN to the WHERE clause

city

city_id	city	country_id	last_update
1	A Corua (La Corua)	87	2021-03-06 15:51:49
2	Abha	82	2021-03-06 15:51:49
3	Abu Dhabi	101	2021-03-06 15:51:49
4	Acua	60	2021-03-06 15:51:49
5	Adana	97	2021-03-06 15:51:49
6	Addis Abeba	31	2021-03-06 15:51:49
7	Aden	107	2021-03-06 15:51:49
8	Adoni	44	2021-03-06 15:51:49



country

country_id	country	last_update
1	Afghanistan	2021-03-06 15:51:49
2	Algeria	2021-03-06 15:51:49
3	American Samoa	2021-03-06 15:51:49
4	Angola	2021-03-06 15:51:49
5	Anguilla	2021-03-06 15:51:49
6	Argentina	2021-03-06 15:51:49
7	Armenia	2021-03-06 15:51:49
8	Australia	2021-03-06 15:51:49
9	Austria	2021-03-06 15:51:49

```
SELECT
    city_id, city.city, country.country
FROM
    city
    INNER JOIN country ON city.country_id=country.country_id
;
```

so we add a JOIN to the WHERE clause

# city-and-country

=


city_id	city	country
1	A Corua (La Corua)	Spain
2	Abha	Saudi Arabia
3	Abu Dhabi	United Arab Emirates
4	Acua	Mexico
5	Adana	Turkey
6	Addis Abeba	Ethiopia
7	Aden	Yemen
8	Adoni	India

```
SELECT
    city_id, city.city, country.country
FROM
    city
    INNER JOIN country ON city.country_id=country.country_id
;
```

```
SELECT ... FROM a INNER JOIN b ON ...;
```

- ```
SELECT a.city, b.country  
FROM city a  
    INNER JOIN country b ON a.country_id=b.country_id  
; -- output a table with city-country names
```
- ```
SELECT f.title, f.length, l.name  
FROM film f  
    INNER JOIN language l ON f.language_id=l.language_id  
WHERE rating='R'  
LIMIT 10; -- output a sample of films and the name of the language it is in
```

# FROM a INNER JOIN b



## INNER JOIN




```
SELECT  
    <SELECT LIST>  
FROM    TABLE_A A  
INNER JOIN TABLE_B B  
ON A.KEY = B.KEY
```

KEY	VAL_X	VAL_Y
1	X1	Y1
2	X2	Y2



# LEFT OUTER JOIN



## LEFT JOIN

```
SELECT  
    <SELECT LIST>  
FROM    TABLE_A A  
LEFT JOIN TABLE_B B  
ON A.KEY = B.KEY
```

KEY	VAL_X	VAL_Y
1	X1	Y1
2	X2	Y2
3	X3	NULL





Table 1

A
B
C

Table 2

A
B
D

INNER JOIN: show all matching records in both tables.

A	A
B	B

LEFT JOIN: show all records from left table, and any matching records from right table.

A	A
B	B
C	

RIGHT JOIN: show all records from right table, and any matching records from left table.

A	A
B	B
	D

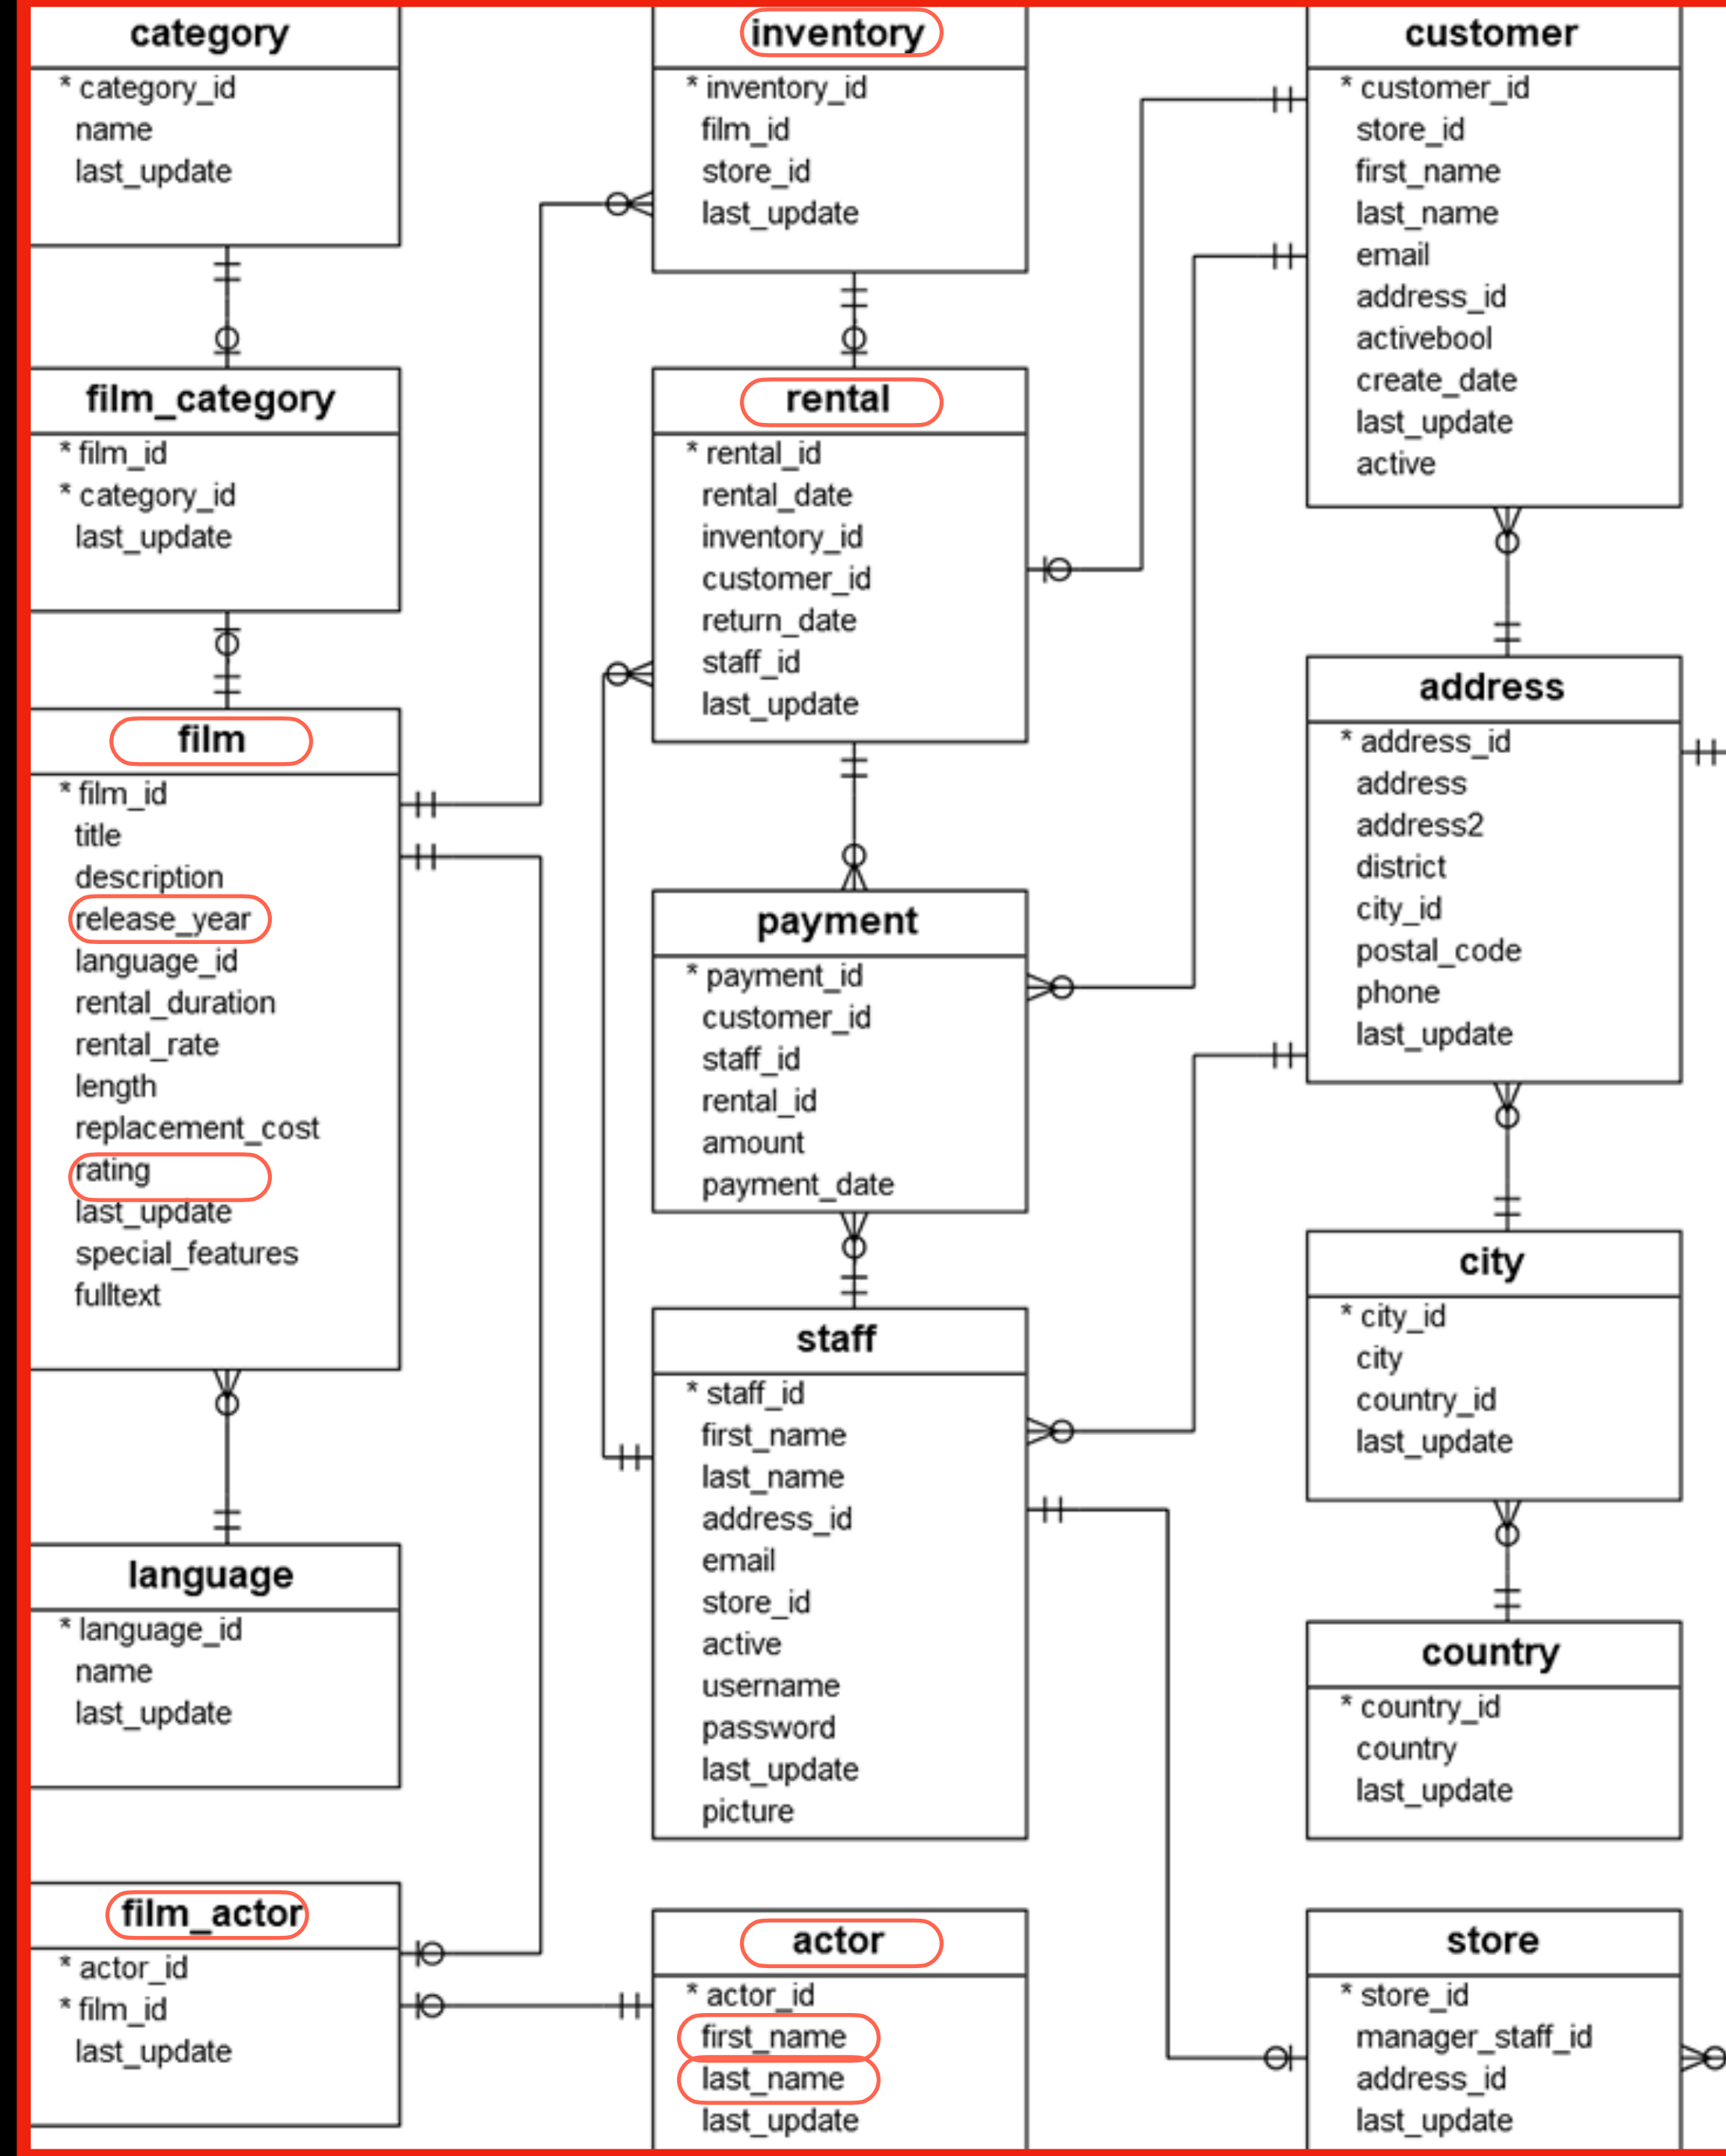
FULL JOIN: show all records from both tables, whether there is a match or not.

A	A
B	B
C	
	D



today's objective:

“which **top 10** actors were rented out the greatest number of times, counting only ‘**R**’ rated films made in **2006**?”



# how do i combine the components of a SELECT?

“which top 10 actors were rented out the greatest number of times, counting only ‘R’ rated films made in 2006?”

- SELECT {columns} FROM {table};
- + LIMIT num
- + WHERE {a\_condition}
- + ORDER BY {columns}
- + GROUP BY {columns}
- + HAVING {a\_condition}
- + [INNER] JOIN {table\_2} ON {col1}={col2}

how the query  
is written

SELECT ...

FROM + JOIN

WHERE ...

GROUP BY ...

HAVING ...

ORDER BY ...

LIMIT ...

how you should  
think about it

FROM + JOIN



WHERE



GROUP BY



HAVING




SELECT



ORDER BY



LIMIT



your turn! compose a query to answer:

“which top 10 actors were rented out the greatest number of times, counting only ‘R’ rated films made in 2006?”









# hint: structure of the solution

```
SELECT
    {}          AS actor_name,
    COUNT({}) AS num_rentals
FROM {}table1
    INNER JOIN {}table2 ON {}join-condition
    INNER JOIN {}table3 ON {}join-condition
    INNER JOIN {}table4 ON {}join-condition
    INNER JOIN {}table5 ON {}join-condition
WHERE {}row condition1
    AND {}row condition2
GROUP BY {}column1
ORDER BY {}column DESC
LIMIT {}num
```

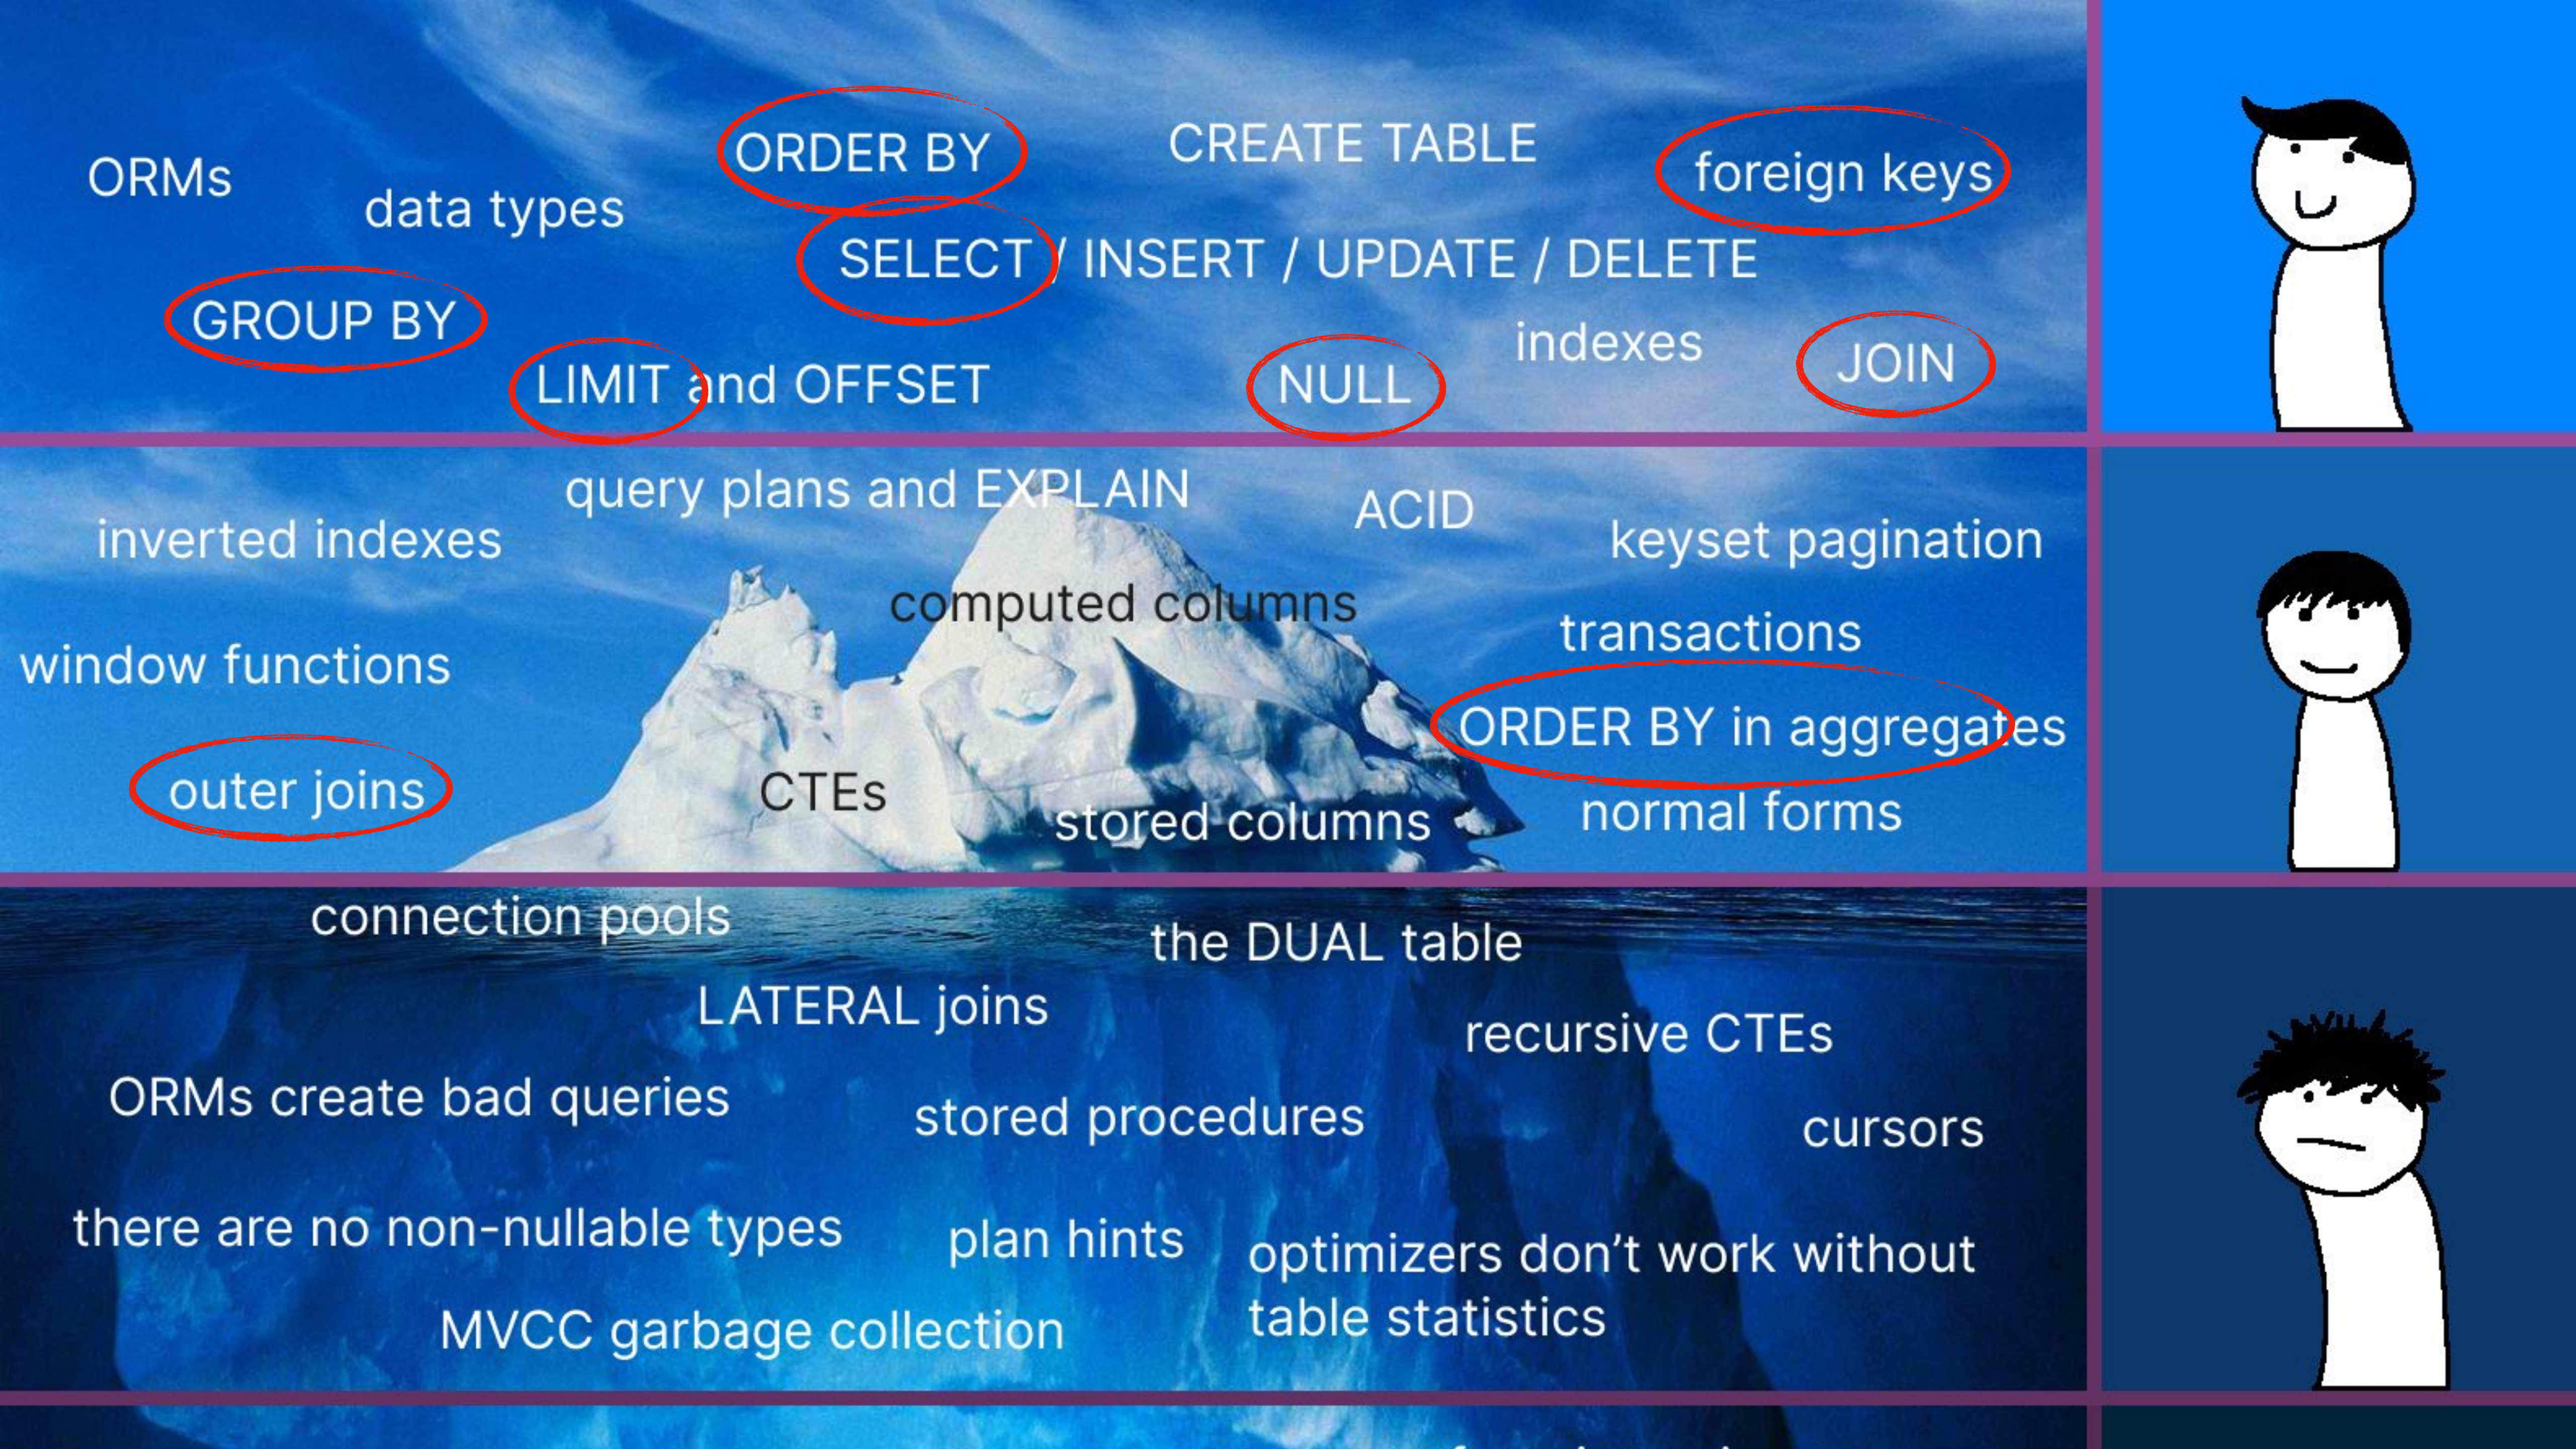






ORMs	data types	ORDER BY	CREATE TABLE	foreign keys	
	GROUP BY	SELECT / INSERT / UPDATE / DELETE	indexes	JOIN	
	LIMIT and OFFSET	NULL			
inverted indexes	query plans and EXPLAIN	ACID	keyset pagination		
computed columns			transactions		
window functions	outer joins	CTEs	ORDER BY in aggregates	normal forms	
stored columns					
connection pools	the DUAL table				
LATERAL joins		recursive CTEs			
ORMs create bad queries	stored procedures	cursors			
there are no non-nullable types	plan hints	optimizers don't work without table statistics			
MVCC garbage collection					
COUNT(*) vs COUNT(1)	isolation levels	generator functions zip when cross joined	sharding		
serializable restarts require retry loops on all statements	zigzag join	phantom reads	triggers	MERGE	
grouping sets, cube, rollup	write skew	partial indexes			
denormalization	SELECT FOR UPDATE	NULLs in CHECK constraints are truthy			
transaction contention	sargability	timestamptz doesn't store a timezone	star schemas		
ascending key problem	ambiguous network errors	utf8mb4			
cost models don't reflect reality	'null':jsonb IS NULL = false	TPCC requires wait times			
DEFERRABLE INITIALLY IMMEDIATE		causal reverse			
EXPLAIN approximates SELECT COUNT(*)	MATCH PARTIAL foreign keys				
vectorized doesn't mean SIMD	NULLs are equal in DISTINCT but inequal in UNIQUE	volcano model			
join ordering is NP hard	database cracking	WCOJ			
learned indexes	XTID exhaustion				
the halloween problem	dee and dum	SERIAL is non-transactional			
fsyncgate	allballs	NULL	every sql operator is actually a join		





ORMs

data types

ORDER BY

CREATE TABLE

foreign keys

GROUP BY

SELECT / INSERT / UPDATE / DELETE

LIMIT and OFFSET

NULL

indexes

JOIN

query plans and EXPLAIN

ACID

keyset pagination

inverted indexes

computed columns

transactions

window functions

ORDER BY in aggregates

normal forms

outer joins

CTEs

stored columns

connection pools

the DUAL table

LATERAL joins

recursive CTEs

ORMs create bad queries

stored procedures

cursors

there are no non-nullable types

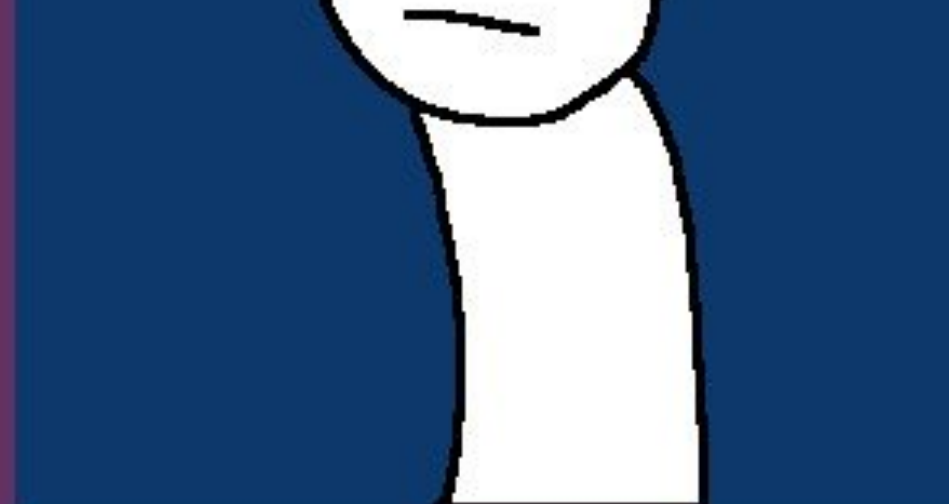
plan hints

optimizers don't work without  
table statistics

MVCC garbage collection



there are no non-nullable types  
plan hints  
optimizers don't work without table statistics  
MVCC garbage collection



COUNT(\*) vs COUNT(1)  
isolation levels  
generator functions zip when cross joined  
sharding  
serializable restarts require retry loops on all statements  
zigzag join  
phantom reads  
triggers  
MERGE  
grouping sets, cube, rollup  
write skew  
partial indexes



denormalization  
SELECT FOR UPDATE  
NULLs in CHECK constraints are truthy  
star schemas  
transaction contention  
sargability  
timestamptz doesn't store a timezone  
utf8mb4  
ascending key problem  
ambiguous network errors



cost models don't reflect reality  
'null'::jsonb IS NULL = false  
TPCC requires wait times  
DEFERRABLE INITIALLY IMMEDIATE





cost models don't  
reflect reality

EXPLAIN approximates  
SELECT COUNT(\*)

'null'::jsonb IS NULL = false  
DEFERRABLE INITIALLY IMMEDIATE

MATCH PARTIAL foreign keys

TPCC requires wait times

causal reverse



vectorized doesn't  
mean SIMD

NULLs are equal in DISTINCT  
but inequal in UNIQUE

volcano model

join ordering is NP hard

database cracking

WCOJ

learned indexes

XTID exhaustion



the halloween problem

dee and dum

SERIAL is non-transactional

allballs

fsyncgate

NULL

every sql operator is  
actually a join





# further learning

- refresher:  
<https://www.youtube.com/watch?v=kbKty5ZVKMY>
- pandas experts note:  
<https://www.youtube.com/watch?v=fmrmwFPMMaM>
- more discussion:  
<https://www.youtube.com/watch?v=OV6Mh2JI9zQ>
- deeper learning:  
<https://app.datacamp.com/learn/career-tracks/data-analyst-in-sql>
- two week free course online starting 2023-02-20:  
<https://corise.com/course/sql-crash-course>

